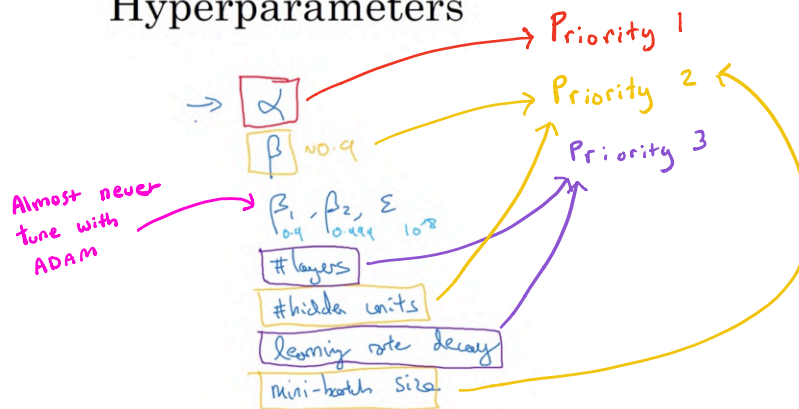
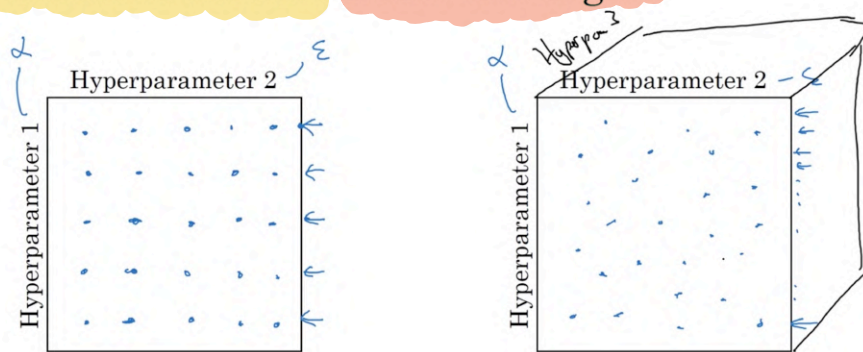


## Hyperparameters



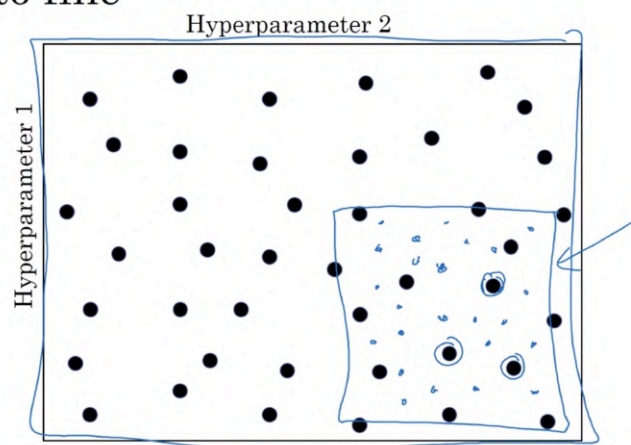
Andrew Ng

Try random values: Don't use a grid



Andrew Ng

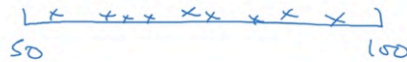
Coarse to fine



Andrew Ng

## Picking hyperparameters at random

$$\rightarrow n^{T2T} = 50, \dots, 100$$



$$\rightarrow \# \text{layers } L: 2 - 4$$

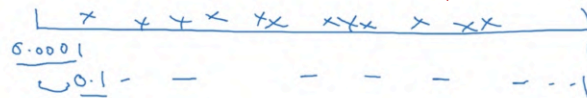
$$2, 3, 4$$

OR

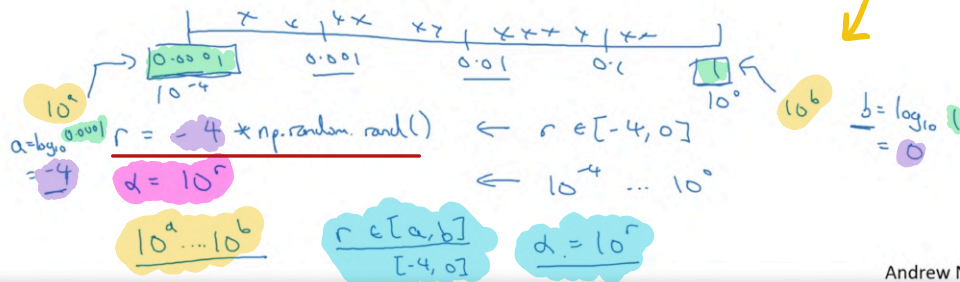
Andrew Ng

## Appropriate scale for hyperparameters

$$\alpha = 0.0001, \dots, 1$$



Uniform  
log scale



Andrew Ng

so  $\alpha = 10^r$  where  $r \in [-4, 0]$

## Hyperparameters for exponentially weighted averages

$$\beta = 0.9 \dots 0.999$$

$$\downarrow \quad \quad \downarrow$$

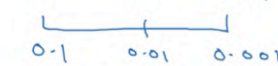
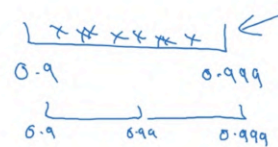
$$10 \quad \quad 1000$$

$$1 - \beta = 0.1 \dots 0.001$$

$$\beta: 0.9000 \rightarrow 0.9005 \} \sim 10$$

$$\beta: 0.999 \rightarrow 0.9995 \} \sim 1000$$

$$\frac{1}{1 - \beta}$$



$$10^{-1} \quad \quad 10^{-3}$$

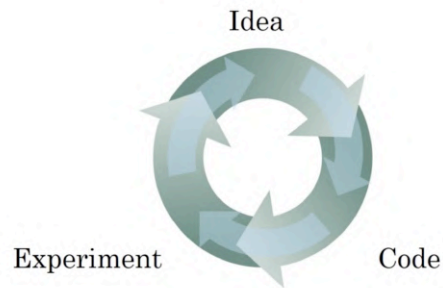
$$r \in [-3, -1]$$

$$1 - \beta = 10^r$$

$$\beta = 1 - 10^r$$

Andrew Ng

## Re-test hyperparameters occasionally

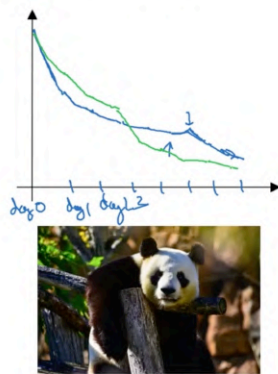


- NLP, Vision, Speech,  
Ads, logistics, ....

- Intuitions do get stale.  
Re-evaluate occasionally.

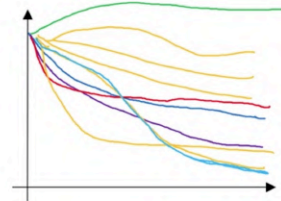
Andrew Ng

### Babysitting one model



Panda ←

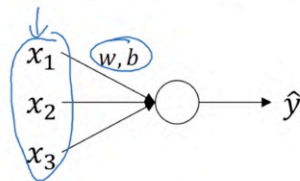
### Training many models in parallel



Caviar ←

Andrew Ng

## Normalizing inputs to speed up learning



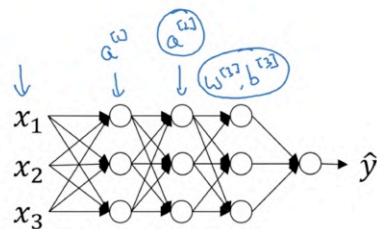
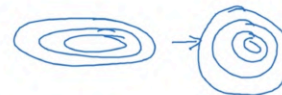
$$\mu = \frac{1}{n} \sum_i x^{(i)}$$

$$X = X - \mu$$

$$\sigma^2 = \frac{1}{n} \sum_i x^{(i)2}$$

$$X = X / \sigma^2$$

← element-wise

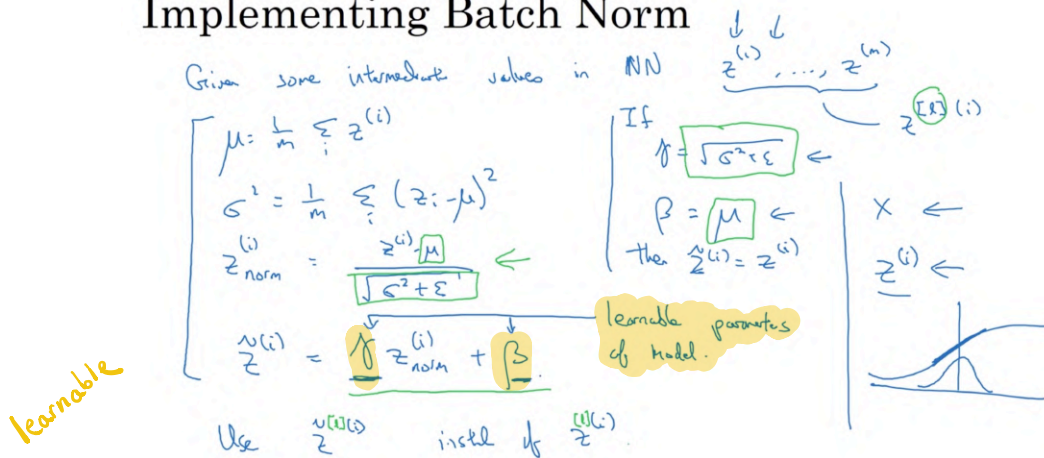


Can we normalize  $\frac{a^{[2]}}{w^{[2]}, b^{[2]}}$  so as to train faster

Normalize  $\frac{z^{[2]}}{\uparrow}$

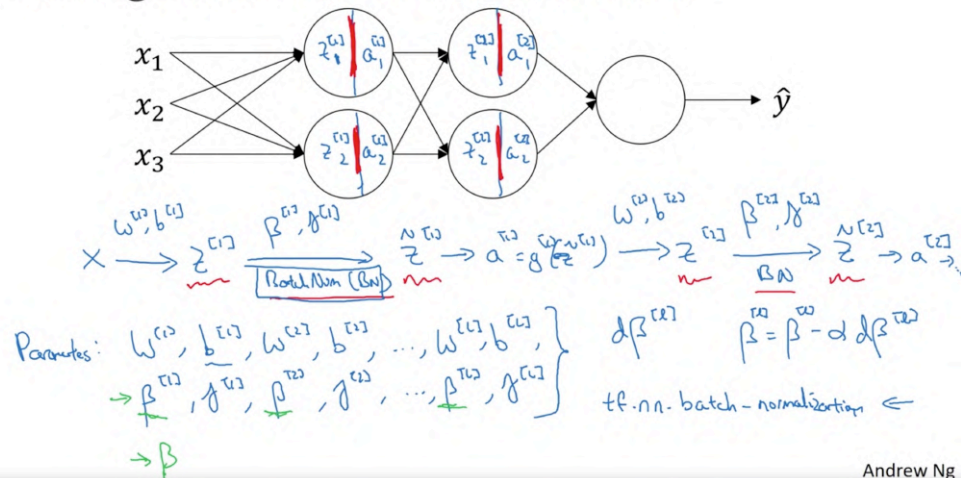
Andrew Ng

## Implementing Batch Norm



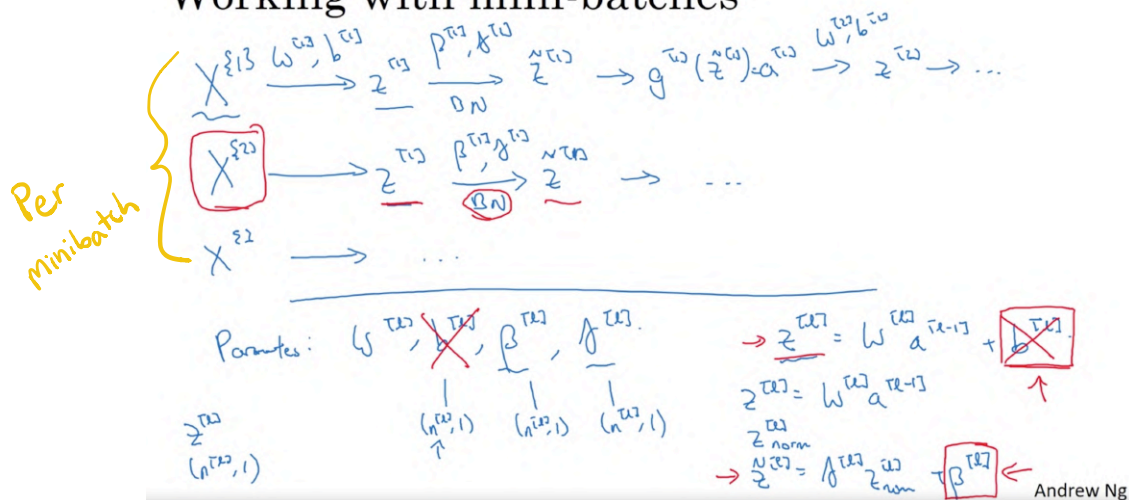
Andrew Ng

## Adding Batch Norm to a network



Andrew Ng

## Working with mini-batches





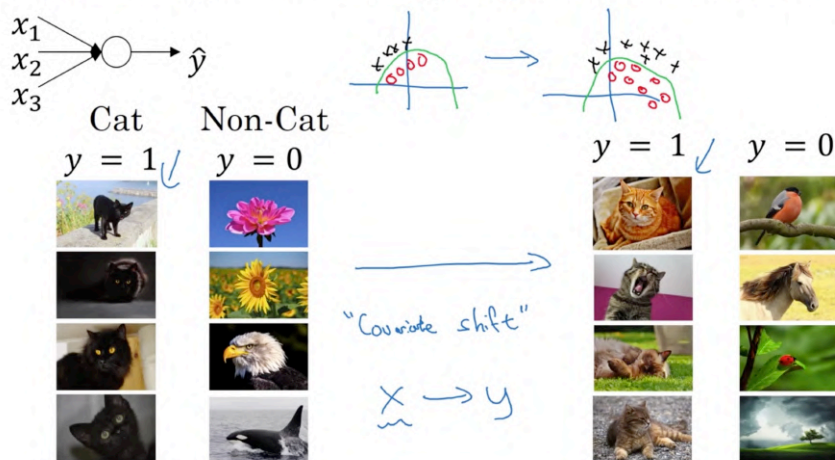
## Implementing gradient descent

for  $t = 1 \dots \text{num MiniBatches}$   
 Compute forward pass on  $X^{\text{test}}$ .  
 In each hidden layer, use BN to replace  $z^{\text{test}}$  with  $\hat{z}^{\text{test}}$ .  
 Use backprop to compute  $\frac{dL}{dW^{\text{test}}}$ ,  $\frac{dL}{d\beta^{\text{test}}}$ ,  $\frac{dL}{d\gamma^{\text{test}}}$ .  
 Update parameters  $\left. \begin{aligned} W^{\text{test}} &:= W^{\text{test}} - \alpha \frac{dL}{dW^{\text{test}}} \\ \beta^{\text{test}} &:= \beta^{\text{test}} - \alpha \frac{dL}{d\beta^{\text{test}}} \\ \gamma^{\text{test}} &:= \dots \end{aligned} \right\} \leftarrow$

Works w/ momentum, RMSprop, Adam.

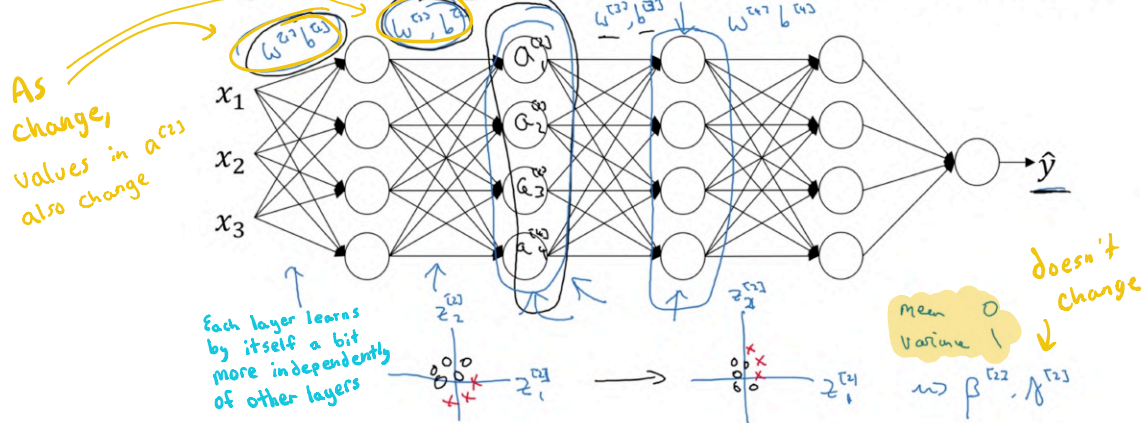
Andrew Ng

## Learning on shifting input distribution



Andrew Ng

## Why this is a problem with neural networks?



Andrew Ng

## Batch Norm as regularization

- Each mini-batch is scaled by the mean/variance computed on just that mini-batch.
- This adds some noise to the values  $z^{[l]}$  within that minibatch. So similar to dropout, it adds some noise to each hidden layer's activations.
- This has a slight regularization effect.

batch norm with dropout for greater regularization effect of batch norm

I wouldn't really use batch norm as a regularizer, that's really not the intent of batch norm, but sometimes it has this extra intended or unintended effect on your learning algorithm. But, really, don't turn to batch norm as a regularization. Use it as a way to normalize your hidden units activations and therefore speed up learning. And I think the regularization is an almost unintended side effect

Andrew Ng

## Batch Norm at test time

$$\begin{aligned} \mu &= \frac{1}{m} \sum_i z^{(i)} \\ \sigma^2 &= \frac{1}{m} \sum_i (z^{(i)} - \mu)^2 \\ z_{\text{norm}}^{(i)} &= \frac{z^{(i)} - \mu}{\sqrt{\sigma^2 + \epsilon}} \\ \tilde{z}^{(i)} &= \gamma z_{\text{norm}}^{(i)} + \beta \end{aligned}$$

$\mu, \sigma^2$ : estimate using exponentially weighted average (across mini-batches).

$x^{[2]}, x^{[3]}, \dots$

$\mu^{[2]}, \mu^{[3]}, \dots$

$\sigma_1^2, \sigma_2^2, \sigma_3^2, \dots$

$\tilde{z}_{\text{norm}} = \frac{z - \mu}{\sqrt{\sigma^2 + \epsilon}}$

$\tilde{z} = \gamma \tilde{z}_{\text{norm}} + \beta$

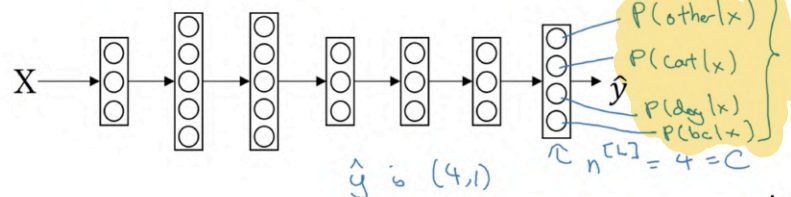
Andrew Ng

## Recognizing cats, dogs, and baby chicks, other



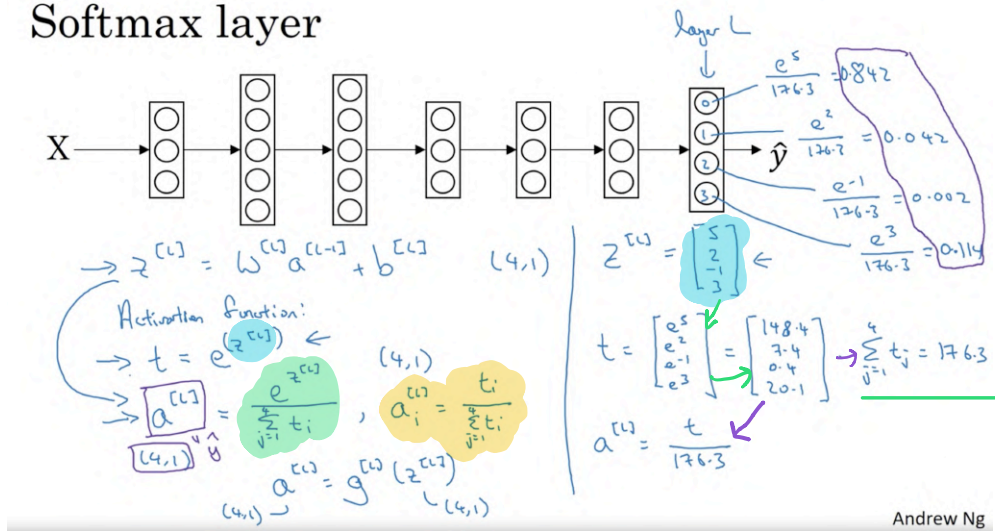
3 1 2 0 3 2 0 1

$C = \text{\#classes} = 4$  (0, ..., 3)

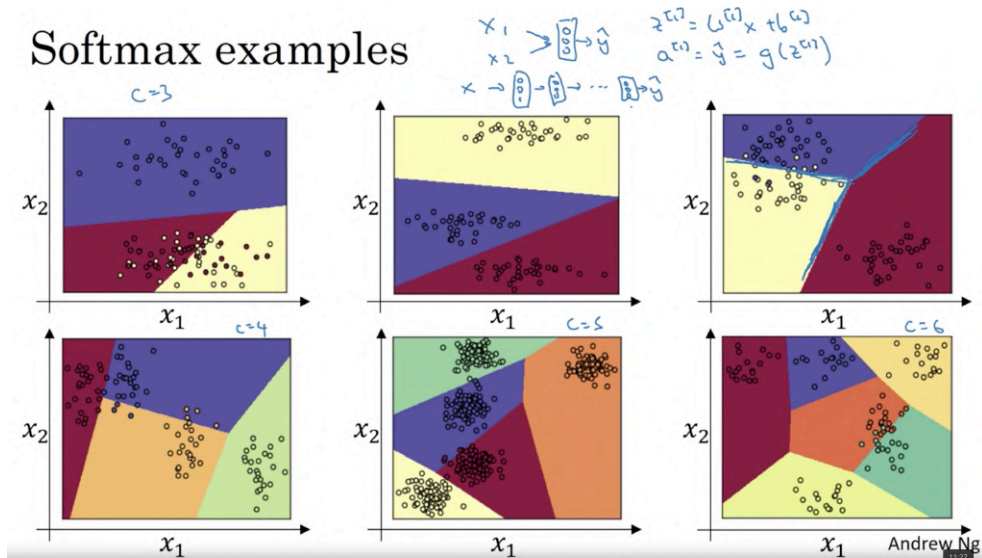


Andrew Ng

## Softmax layer



## Softmax examples



## Understanding softmax

Handwritten calculation for  $C=4$ :

$$z^{(L)} = \begin{bmatrix} 5 \\ 2 \\ -1 \\ 3 \end{bmatrix} \quad t = \begin{bmatrix} e^5 \\ e^2 \\ e^{-1} \\ e^3 \end{bmatrix}$$

Softmax function:

$$a^{(L)} = g^{(L)}(z^{(L)}) = \begin{bmatrix} \frac{e^5}{e^5 + e^2 + e^{-1} + e^3} \\ \frac{e^2}{e^5 + e^2 + e^{-1} + e^3} \\ \frac{e^{-1}}{e^5 + e^2 + e^{-1} + e^3} \\ \frac{e^3}{e^5 + e^2 + e^{-1} + e^3} \end{bmatrix} = \begin{bmatrix} 0.842 \\ 0.042 \\ 0.002 \\ 0.114 \end{bmatrix}$$

Handwritten notes:

- $C=4$
- $g^{(L)}(\cdot)$
- "hard max" (indicated by a box around the highest value 0.842)

Softmax regression generalizes logistic regression to  $C$  classes.

If  $C=2$ , softmax reduces to logistic regression.  $a^{(L)} = \begin{bmatrix} 0.842 \\ 0.158 \end{bmatrix}$

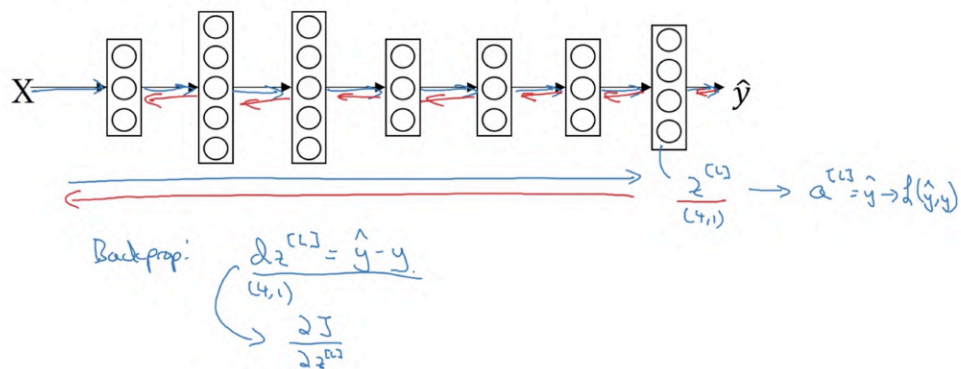


## Loss function

$(4,1)$   
 $y = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$  - cat  $y_2 = 1$   
 $y_1 = y_2 = y_3 = y_4 = 0$   
 $(4,1)$   
 $\hat{y} = \begin{bmatrix} 0.3 \\ 0.2 \\ 0.1 \\ 0.4 \end{bmatrix}$   
 $C = 4$   
 $\mathcal{L}(\hat{y}, y) = - \sum_{j=1}^C y_j \log \hat{y}_j$   
 $\mathcal{J}(w^{(1)}, b^{(1)}, \dots) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$   
 $-y_2 \log \hat{y}_2 = -\log \hat{y}_2 \rightarrow \text{make } \hat{y}_2 \text{ big.}$   
 $Y = [y^{(1)} \ y^{(2)} \ \dots \ y^{(n)}]$   
 $\hat{Y} = [\hat{y}^{(1)} \ \dots \ \hat{y}^{(n)}]$   
 $= \begin{bmatrix} 0 & 0 & 1 & 0 & \dots \\ 1 & 0 & 0 & 0 & \dots \\ 0 & 0 & 0 & 0 & \dots \\ 0 & 0 & 0 & 0 & \dots \end{bmatrix}$   
 $(4, m)$   
 $= \begin{bmatrix} 0.3 & \dots \\ 0.2 & \dots \\ 0.1 & \dots \\ 0.4 & \dots \end{bmatrix}$   
 $(4, m)$

Andrew Ng

## Gradient descent with softmax



Andrew Ng