Assignment 2: Transformer Summarizer

Welcome to the second assignment of course 4. In this assignment you will explore summarization using the transformer model. Yes, you will implement the transformer decoder from scratch, but we will slowly walk you through it. There are many hints in this notebook so feel free to use them as needed.

Outline

- Introduction
- Part 1: Importing the dataset
 - 1.1 Encode & Decode helper functions
 - 1.2 Defining parameters
 - 1.3 Exploring the data
- Part 2: Summarization with transformer
 - 2.1 Dot product attention
 - Exercise 01
 - 2.2 Causal Attention
 - Exercise 02
 - 2.3 Transformer decoder block
 - Exercise 03
 - 2.4 Transformer Language model
 - Exercise 04
- Part 3: Training
 - 3.1 Training the model
 - Exercise 05
- Part 4: Evaluation
 - 4.1 Loading in a trained model
- Part 5: Testing with your own input
 - Exercise 6
 - 5.1 Greedy decoding
 - Exercise 07

Introduction

Summarization is an important task in natural language processing and could be useful for a consumer enterprise. For example, bots can be used to scrape articles, summarize them, and then you can use sentiment analysis to identify the sentiment about certain stocks. Anyways who wants to read an article or a long email today, when you can build a transformer to summarize text for you. Let's get started, by completing this assignment you will learn to:

- Use built-in functions to preprocess your data
- Implement DotProductAttention
- Implement Causal Attention
- Understand how attention works
- · Build the transformer model
- Evaluate your model
- Summarize an article

As you can tell, this model is slightly different than the ones you have already implemented. This is heavily based on attention and does not rely on sequences, which allows for parallel computing.

```
In [24]:
```

```
import sys
import os

import numpy as np

import textwrap
wrapper = textwrap.TextWrapper(width=70)

import trax
```

```
from trax import layers as tl
from trax.fastmath import numpy as jnp

# to print the entire np array
np.set_printoptions(threshold=sys.maxsize)
```

Part 1: Importing the dataset

Trax makes it easy to work with Tensorflow's datasets:

In [25]:

1.1 Tokenize & Detokenize helper functions

Just like in the previous assignment, the cell above loads in the encoder for you. Given any data set, you have to be able to map words to their indices, and indices to their words. The inputs and outputs to your <u>Trax</u> models are usually tensors of numbers where each number corresponds to a word. If you were to process your data manually, you would have to make use of the following:

- word2Ind: a dictionary mapping the word to its index.
- ind2Word: a dictionary mapping the index to its word.
- word2Count: a dictionary mapping the word to the number of times it appears.
- num_words: total number of words that have appeared.

Since you have already implemented these in previous assignments of the specialization, we will provide you with helper functions that will do this for you. Run the cell below to get the following functions:

- tokenize: converts a text sentence to its corresponding token list (i.e. list of indices). Also converts words to subwords.
- detokenize: converts a token list to its corresponding sentence (i.e. string).

In [26]:

1.2 Preprocessing for Language Models: Concatenate It!

This week you will use a language model -- Transformer Decoder -- to solve an input-output problem. As you know, language models only predict the next word, they have no notion of inputs. To create a single input suitable for a language model, we concatenate inputs with targets putting a separator in between. We also need to create a mask -- with 0s at inputs and 1s at targets -- so that the model is not penalized for mis-predicting the article and only focuses on the summary. See the preprocess function below for how this is done.

In [27]:

```
# Special tokens
SEP = 0 # Padding or separator token
EOS = 1 # End of sentence token
# Concatenate tokenized inputs and targets using 0 as separator.
def preprocess(stream):
   for (article, summary) in stream:
       joint = np.array(list(article) + [EOS, SEP] + list(summary) + [EOS])
       mask = [0] * (len(list(article)) + 2) + [1] * (len(list(summary)) + 1) # Accounting for EOS
and SEP
       yield joint, joint, np.array(mask)
# You can combine a few data preprocessing steps into a pipeline like this.
input pipeline = trax.data.Serial(
   # Tokenizes
   trax.data.Tokenize(vocab dir='vocab dir/',
                      vocab file='summarize32k.subword.subwords'),
   # Uses function defined above
   preprocess,
    # Filters out examples longer than 2048
   trax.data.FilterByLength(2048)
# Apply preprocessing to data streams.
train stream = input pipeline(train stream fn())
eval stream = input pipeline(eval stream fn())
train input, train target, train mask = next(train stream)
assert sum((train input - train target)**2) == 0 # They are the same in Language Model (LM).
```

In [28]:

```
# prints mask, 0s on article, 1s on summary
print(f'Single example mask:\n\n {train_mask}')
```

Single example mask:

```
 \  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\  \, 0\ \,
```

In [29]:

```
# prints: [Example][<EOS>][<pad>][Example Summary][<EOS>]
print(f'Single example:\n\n {detokenize(train_input)}')
```

Single example:

President Barack Obama is getting off the island. In a rare move for him, the president planned a break in the middle of his Martha's Vineyard vacation to return to Washington on Sunday night for meetings with Vice President Joe Biden and other advisers on the U.S. military campaign in Iraq and tensions between police and protesters in Ferguson, Missouri. The White House has been cagey about why the president needs to be back in Washington for those discussions. He's received multiple briefings on both issues while on vacation. The White House had also already announced Obama's plans to return to Washington before the U.S. airstrikes in Iraq began and before the shooting of a teen in Ferguson that sparked protests. Protective $% \left(1\right) =\left(1\right) \left(1\right) +\left(1\right) \left(1\right) \left(1\right) +\left(1\right) \left(1\right)$ gesture: Obama walks with daughter Malia Obama to board Air Force One at Cape Cod Coast Guard Air Station in Massachusetts on Sunday . Back home: Obama and Malia are seen at Joint Base Andrews in Washington early Monday . Mysterious: The White House has been cagey about why the president needs to be back in Washington. He is seen here on the South Lawn of the White House with daughter Malia . In good spirits: Despite the early return, The President and First Daughter seemed to be enjoying a joke . Part of the decision to head back to Washington appears aimed at countering criticism that Obama is spending two weeks on a resort island in the midst of so many foreign and domestic crises. Yet those crises turned the first week of Obama's vacation into a working holiday. He made on-camera statements Iraq and the clashes in Ferguson, a St. Louis suburb. He also called foreign leaders to discuss the tensions between Ukraine and Russia, as well as between Israel and Hamas. 'I think it's fair to say there are, of course, ongoing complicated situations in the world, and that's why you've seen the president stay engaged,' White House spokesman Eric Schultz said. Obama returned from his break along with his 16-year-old daughter Mailia, but is scheduled to return to Martha's Vineyard on Tuesday and stay through next weekend. In a first for Obama family summer vacations, neither teenager is spending the entire holiday with her father. Obama left Washington Aug. 9 with his wife, Michelle, daughter Malia, and the family's two Portuguese water dogs. The White House said 13-year-old Sasha would join her parents at a later date for "part of their stay" on this quaint island of shingled homes. But Malia will not be around when her younger sister arrives. The daughters essentially are trading places, and the vacation is boiling down to Obama getting about a week with each one. Malia returned to Washington with her father and is not expected to go back to Martha's Vineyard. The White House said Sasha will join her parents this week, without saying when she will arrive or what kept her away last week, or why Malia left the island. President Barack Obama bike rides with daughter Malia Obama while on vacation with his family on the island of Martha's Vineyard . Obama often draws chuckles from sympathetic parents who understand his complaints about his girls' lack of interest in spending time with him. 'What I'm discovering is that each year, I get more excited about spending time with them. They get a little less excited,' Obama told CNN last year. Even though work has occupied much of Obama's first week on vacation, he still found plenty of time to golf, go to the beach with his family and go out to dinner on the island. He hit the golf course one more time Sunday ahead of his departure, joining two aides and former NBA player Alonzo Mourning for an afternoon round. He then joined wife Michelle for an evening jazz performance featuring singer Rachelle Ferrell. Obama's vacation has also been infused with a dose of politics. He headlined a fundraiser on the island for Democratic Senate candidates and attended a birthday party for Democratic adviser Vernon Jordan's wife, where he spent time with former President Bill Clinton and Hillary Rodham Clinton. That get-together between the former rivals-turned-partners added another complicated dynamic to Obama's vacation. Just as Obama was arriving on Martha's Vineyard, an interview with the former

secretary of state was published in which she levied some of her sharpest criticism of Obama's foreign policy. Clinton later promised she and Obama would 'hug it out' when they saw each other at Jordan's party. No reporters were allowed in, so it's not clear whether there was any hugging, but the White House said the president danced to nearly every song.<EOS><pad>PresidentObama will head back to the White House on Sunday night as tensions rise in Missouri and Iraq . The decision appears aimed at countering criticism that the president was spending two weeks on a resort island in the midst of so many crises

1.3 Batching with bucketing

As in the previous week, we use bucketing to create batches of data.

In [30]:

```
# Bucketing to create batched generators.
# Buckets are defined in terms of boundaries and batch sizes.
# Batch sizes[i] determines the batch size for items with length < boundaries[i]
# So below, we'll take a batch of 16 sentences of length < 128 , 8 of length < 256,
# 4 of length < 512. And so on.
boundaries = [128, 256, 512, 1024]
batch_sizes = [16,
                          4,
                               2, 1]
                    8,
# Create the streams.
train batch stream = trax.data.BucketByLength(
   boundaries, batch sizes) (train stream)
eval batch stream = trax.data.BucketByLength(
   boundaries, batch sizes) (eval stream)
```

In [31]:

```
# Every execution will result in generation of a different article
# Try running this cell multiple times to see how the length of the examples affects the batch siz
input batch, , mask batch = next(train batch stream)
# Shape of the input batch
input batch.shape
```

Out[31]:

(1.1175)

In [32]:

```
# print corresponding integer values
print(input batch[0])
[ 1668 4375 12915 10077 1595
                           592
                                 15 18621
                                            320 21427
                                                            61
                226 1668
                           276 10583 3345
           88
                                            320 15
                                                      205
                                                            7
  320
       32
    5 1698 2572
                   2 1480
                           229
                                     213
                                                 10
                           379
 2572 1248 2058 1782
                      198
                                  49 1151
                                            92
                                                404 1098
                                                           341
                                 18 1537
 2572 1098
                                           403
                                                 379
             2.
                 186 22015 1998
                                                     2.08
                                                            2.8
 1375
      1019
            213
                 1080
                      266 29725
                                           2565
                                                     4182
                                  4
                                       5
                                                 320
 2572
                      4375 127
                                 809
                                       2.8
                                           782 1900
      3898
            213
                 3837
                                                      824
 1782
       9
           715
                 13 410 20688
                                 592
                                     39 11832
                                                824 2854
    6
        78
           691 379 12830 1049
                                 89 1338
                                          320 5421
                                                      213 26246
             2
                 455 3223 22728
                                            379
   84
       889
                                5889
                                     186
                                                727 19064 1779
13862
        4
            213
                1646
                      527
                           101
                                 467
                                     1668
                                            186
                                                 379
                                                      561
                                                           2002
 1324
       320
             28 15324
                       4 10256
                                 17
                                      532
                                            220
                                                 736
                                                      320
                           715
25316
        4
             2 213 2462
                                 39
                                     919 1668 22752
                                                       5
                                                           281
  225
       446
            28 984
                       3 9175 6051
                                           246 1019
                                                     846
24408
        4 4706 239
                      11 1668 4375 12915 10077
                                                127
                                                     824 4979
                           61
                                320 32
  285
        22
            229 21427
                                           88
                                                 226
                       16
                                                     1668
                                                           276
10583 3345
            320
                213
                       205
                             7
                                  5
                                     1698
                                           2572
                                                 2
                                                     1480
                           3 2572 1248
           184
   43
      213
                  10 59
                                           2058
                                                379 23736
                           213 1668
        2 10077
                23 24462
                                     676
                                           527 1347 8048
```

_ ~	-						٠. ٥	~ <u>.</u> .			~ = v
835		132				786					
919 213	527 72		32 39		53 281	446 65	3 446			500 2	2 213
15324		10209	834		186		103			869	
213	750		413	379	1838	132	213		80	5.4	7.4
117	442	1385			107		662				9
1624 10127	132 71		1660 184			28	18304			199 1356	
2577	88	226	184 14160	310	2	125	527			19357	
17859	21	2	18	14513	918	2251	213			220	
2 691	186	213		4150 824		1549			226 59	39	
26156	213 4	704 23				3 691				14	24684
186	213			229		1260				320	
1019	213	310	3			132				18994	
281 285	53 62	10 1422		750		1767 213		1838		26651 35	4217 638
9229		13341	67			1155				3	
24299	20	605	527	213	750	4217	7	5	4823		62
273	1629					213				445	
273 9229	1629 62	4925 107	213 320			31 320				378	
2193	213	266	320	616	310	1838			12703		431
628	1779	344				2572					
428 1094	171 711	41	196	1151	1233	278 2	5309	2/021	7166	428 3681	
13365	16	246	213	9815	13277	361	428	3	9	1155	40
3640	2935	320	377	28	8010	1036	527	213	472	2407	132
15	2418	320	1608	285	62	18 14715	1325	213	676	527	
268 617	3283 428	320 691	3251 1657			3466					955
278	3	4217	8703	246	809	213	220	2425	102	3666	
6559						10077					
8480 2	747 1668	12099				231					
	10583	3345	25	1048	320	213 5421	19064	285	1435	17889	16
	18304	527	310	186	1144	5709	213	184	10	59	3
	918	379	1312	19	569	132 10583	4217	7	5	2418	320 70
1608 579	1353 638	1767 7049	1019	26698	276 186	213	1668	4375	40	2572 148	
78	213	1155				9229				10583	1367
229	1048	809	370	527	208	1881	320	399	24684		
2846	18246	246				186 4217					
1668											
213											4 4217
15194	17 8676					892					
2 1912		2511	378	4606	1631	3345 3	303	10630	8962	117	5773
20	80	9958	13589	12650	2	28	17376	1779	2667	2572	513
	21749		2	14651	213	3837	4375	7	5	20135	412
16399 2900	1133 527	6053 213	1359	3898	22	26246 793	902 213	27872	391	1782	200
89	280	378				26319					
527	213 213	224	5614	320	213	224 28	864	7151	1435	892	662
527 726	213 1106		1782 397	56 51	229	28 229	1424	1106	320	19	28 44
25510	283					26156				13589	
127		326	1435	506	101	2	141	1144		467	
207 2002		165 10256				207					
13804		285	103	229	28	78 117	22531	4505	1700	527	213
2572	3898	487	3	244	10077	7	5	797	6107	17946	21
592	285	3345	62			10943					
1248 1779	378 1086	4606 24792	1631	2002	9958 1248	13589 10940	12650	43 554	5754 412	160	2 527
28	707					527					
15	2575		213	1359	809	213	2572	1782	337	379	97
7230 13542	1083 379		320	2572	2	41 141	358	7	26	662	2685
340	3898			10		0					
824	4979		22	39	14037	4	61	320	32	88	226
1668		10583				2572					
20135 931	39 320					281 4					
23	1065	4217	1547	450	320	2733	213	276	10583	320	213
2572	3.5	4217	8283	25370	16	16346	27439	6774	1628	27	1668

-			1-1	0200	200,0		10010	- ,	0 / / 1	1020	- 1	
26	792	2009	127	213	3837	5365	229	19	19677	114	2331	2685
	213	2572	2	22	141	2976	320	385	1716	2104	11	

Things to notice:

- First we see the corresponding values of the words.
- The first 1, which represents the <EOS> tag of the article.
- Followed by a 0, which represents a <pad> tag.
- After the first 0 (<pad> tag) the corresponding values are of the words that are used for the summary of the article.
- The second 1 represents the <EOS> tag for the summary.
- All the trailing 0s represent <pad> tags which are appended to maintain consistent length (If you don't see them then it would mean it is already of max length)

In [33]:

```
# print the article and its summary
print('Article:\n\n', detokenize(input_batch[0]))
```

Article:

Texas Governor Rick Perry announced today his intentions to deploy up to 1,000 Texas National Guard troops to his state's southern border, which is also the U.S. border with Mexico. 'There . can be no national security without border security, and Texans have paid too . high \boldsymbol{a} price for the federal government's failure to secure our border,' the Republican Governor said at a news conference this afternoon. 'The action I am ordering today will tackle this crisis head-on by . multiplying our efforts to combat the cartel activity, human traffickers and . individual criminals who threaten the safety of people across Texas and . America.' According to a memo leaked late last night to The Monitor, the executive action will cost Texas taxpayers \$12 million a month. Scroll down for video . Done waiting around: Texas Governor Rick Perry said this afternoon that he is deploying up to 1,000 Texas National Guard troops to the state's southern border, which is also the U.S. border with Mexico . Already, Perry has instructed the Texas Department of Public Safety to increase personnel in the Rio Grande River Valley area at a weekly cost of \$1.3 million. Added together, the two measures will cost \$5 million a week, the memo reportedly states, and 'it is not clear where the money will come . from in the budget' other than 'non critical' areas like health care or transportation. The rise in border protection measures follows a surge of Central American children streaming into the U.S. from Mexico. More than 57,000 immigrant children, many of whom are unaccompanied, have illegally entered the country since last year, and the government estimates that approximately 90,000 will arrive by the close of this year. U.S. Border Patrol has been overloaded by the deluge, and the federal government is quickly running out of money to care for the children. Congress is in the process of reviewing a \$3.7 billion emergency funding request from President Barack Obama that would appropriate additional money to the agencies involved, but House Republicans remain skeptical of the president's plan. Roughly half of the money Obama's asking for would go toward providing humanitarian aid to the children while relatively little would go toward returning the them to their home countries. Furthermore, Republicans would like to see changes to a 2008 trafficking law that requires the government to give children from non-contiguous countries who show up at the border health screenings and due process before they can be sent home. The judicial process often takes months, and even years, clogging up courts and slowing down the repatriation process. The president had initially planned to include a revised version of the 2008 legislation in his request to Congress that would have allowed the Department of Homeland Security to exercise the 'discretion' to bypass the current process by giving children the option to voluntarily return home. Obama backed down at the last minute after receiving negative feedback from Democratic lawmakers. Perry held a news conference with Attorney General Greg Abbott, right, this afternoon in Austin, Texas, to formally announce the deployment. Perry said the National Guard troops were needed to combat criminals that are exploiting a surge of children and families entering the U.S. illegally . Also not included in Obama's request to Congress was funding for a National Guard deployment to the border - something House Speaker John Boehner and the Texas Governor had both called on the president to do. Republicans

say a National Guard presence is needed at areas of high crime to help Border Patrol agents crack down on smugglers and drug cartels. In a face to face meeting with Obama when the president came to Texas two weeks ago Perry again asked the president to deploy the National Guard through a federally funded statue but Obama resisted. Perry is now taking matters into his own hands, sending his own set of troops down to the Rio Grande Valley to aid law enforcement officials. State Senator Juan 'Chuy' Hinojosa, a Democrat who represents border town McAllen, criticized the Republican Governor's deployment as unnecessary. '[The cartels] are taking advantage of the situation,' he told the Monitor. 'But our local law enforcement from the sheriff's offices of the different counties to the different police departments are taking care of the situation. 'This is a civil matter, not a military matter. What we need is more resources to hire more deputies, hire more Border Patrol, 'Hinojosa said. 'These are young people, just families coming across. They're not armed. They're not carrying weapons.' The leaked memo on the National Guard deployment specifically denies that it is a 'militarization of the border,' however. And Perry's office reiterated today that troops would 'work seamlessly and side by side with law enforcement officials.' Hinojosa also accused Perry, who recently toured the border with Sean Hannity as part of a special for Fox News, of being insincere in his concern about the situation at the border. 'All . these politicians coming down to border, they don't care about solving . the problem, they just want to make a political point, 'he said. <EOS> <pad>TexasGovernor Rick Perry announced this afternoon that he will dispatch up to 1,000 Texas National Guard troops to the border . The deployment will cost Texastaxpayers \$12 million a month, according to a leaked memo . Perry has asked Obama multiple times to send the National Guard to the border but Obama keeps refusing . A Texas lawmaker said the Republican governor is not sincerely concerned about the border, he just wants to play politics .<EOS>

You can see that the data has the following structure:

```
• [Article] -> <EOS> -> <pad> -> [Article Summary] -> <EOS> -> (possibly) multiple <pad>
```

The loss is taken only on the summary using cross_entropy as loss function.

Part 2: Summarization with transformer

Now that we have given you the data generator and have handled the preprocessing for you, it is time for you to build your own model. We saved you some time because we know you have already preprocessed data before in this specialization, so we would rather you spend your time doing the next steps.

You will be implementing the attention from scratch and then using it in your transformer model. Concretely, you will understand how attention works, how you use it to connect the encoder and the decoder.

2.1 Dot product attention

Now you will implement dot product attention which takes in a query, key, value, and a mask. It returns the output.

Here are some helper functions that will help you create tensors and display useful information:

- create tensor creates a jax numpy array from a list of lists.
- display tensor prints out the shape and the actual tensor.

In [34]:

```
def create_tensor(t):
    """Create tensor from list of lists"""
    return jnp.array(t)

def display_tensor(t, name):
    """Display shape and tensor"""
    print(f'{name} shape: {t.shape}\n')
    print(f'{t}\n')
```

Before implementing it yourself, you can play around with a toy example of dot product attention without the softmax operation. Technically it would not be dot product attention without the softmax but this is done to avoid giving away too much of the answer and the idea is to display these tensors to give you a sense of how they look like.

The formula for attention is this one:

 $\$ \text { Attention }(Q, K, V)=\operatorname{\softmax}\\left(\frac{Q K^{T}}{\sqrt{d_{k}}}+{M}\right) \tag{1}\ \$\$

\$d_{k}\$ stands for the dimension of queries and keys.

The query, key, value and mask vectors are provided for this example.

Notice that the masking is done using very negative values that will yield a similar effect to using \$-\infty \$.

```
In [35]:
q = create_tensor([[1, 0, 0], [0, 1, 0]])
display_tensor(q, 'query')
k = create tensor([[1, 2, 3], [4, 5, 6]])
display_tensor(k, 'key')
v = create_tensor([[0, 1, 0], [1, 0, 1]])
display tensor(v, 'value')
m = create_tensor([[0, 0], [-1e9, 0]])
display_tensor(m, 'mask')
query shape: (2, 3)
[[1 0 0]
 [0 1 0]]
key shape: (2, 3)
[[1 2 3]
 [4 5 6]]
value shape: (2, 3)
[[0 1 0]
[1 0 1]]
mask shape: (2, 2)
[[ 0.e+00 0.e+00]
 [-1.e+09 0.e+00]]
```

Expected Output:

```
query shape: (2, 3)
  [[1 0 0]
   [0 1 0]]
  key shape: (2, 3)
   [[1 2 3]
   [4 5 6]]
  value shape: (2, 3)
   [[0 1 0]
   [1 0 1]]
  mask shape: (2, 2)
  [[ 0.e+00 0.e+00]
[-1.e+09 0.e+00]]
```

```
In [13]:
q dot k = q @ k.T / jnp.sqrt(3)
display tensor(q dot k, 'query dot key')
query dot key shape: (2, 2)
[[0.57735026 2.309401
 [1.1547005 2.8867514]]
Expected Output:
   query dot key shape: (2, 2)
  [[0.57735026 2.309401 ]
[1.1547005 2.8867514]]
In [38]:
masked = q dot k + m
display tensor(masked, 'masked query dot key')
masked query dot key shape: (2, 2)
[[ 5.7735026e-01 2.3094010e+00]
 [-1.0000000e+09 2.8867514e+00]]
Expected Output:
   masked query dot key shape: (2, 2)
   [[ 5.7735026e-01 2.3094010e+00]
 [-1.0000000e+09 2.8867514e+00]]
In [39]:
display_tensor(masked @ v, 'masked query dot key dot value')
masked query dot key dot value shape: (2, 3)
[[ 2.3094010e+00 5.7735026e-01 2.3094010e+00]
 [ 2.8867514e+00 -1.0000000e+09 2.8867514e+00]]
Expected Output:
   masked query dot key dot value shape: (2, 3)
   [[ 2.3094010e+00 5.7735026e-01 2.3094010e+00]
[ 2.8867514e+00 -1.0000000e+09 2.8867514e+00]]
In order to use the previous dummy tensors to test some of the graded functions, a batch dimension should be added to them so they
mimic the shape of real-life examples. The mask is also replaced by a version of it that resembles the one that is used by trax:
In [40]:
```

q_with_batch = q[None,:]
display_tensor(q_with_batch, 'query with batch dim')
k_with_batch = k[None,:]
display_tensor(k_with_batch, 'key with batch dim')
v_with_batch = v[None,:]
display_tensor(v_with_batch, 'value with batch dim')
m_bool = create_tensor([[True, True], [False, True]])

```
query with batch dim shape: (1, 2, 3)

[[[1 0 0]
      [0 1 0]]]]

key with batch dim shape: (1, 2, 3)

[[[1 2 3]
      [4 5 6]]]

value with batch dim shape: (1, 2, 3)

[[[0 1 0]
      [1 0 1]]]

boolean mask shape: (2, 2)

[[ True True]
      [False True]]
```

Exercise 01

Instructions: Implement the dot product attention. Concretely, implement the following equation $\$ \text { Attention }(Q, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Attention}(Q, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement the following equation $\$ \text{Normalize}(A, K, V)=\product attention. Concretely, implement attention atte

\$Q\$ - query, \$K\$ - key, \$V\$ - values, \$M\$ - mask, \${d k}\$ - depth/dimension of the queries and keys (used for scaling down)

You can implement this formula either by trax numpy (trax.math.numpy) or regular numpy but it is recommended to use jnp.

Something to take into consideration is that within trax, the masks are tensors of True/False values not 0's and \$-\infty\$ as in the previous example. Within the graded function don't think of applying the mask by summing up matrices, instead use jnp.where() and treat the mask as a tensor of boolean values with False for values that need to be masked and True for the ones that don't.

Also take into account that the real tensors are far more complex than the toy ones you just played with. Because of this avoid using shortened operations such as @ for dot product or .T for transposing. Use <code>jnp.matmul()</code> and <code>jnp.swapaxes()</code> instead.

This is the self-attention block for the transformer decoder. Good luck!

```
In [57]:
```

```
# GRADED FUNCTION: DotProductAttention
def DotProductAttention(query, key, value, mask):
    """Dot product self-attention.
    Args:
        query (jax.interpreters.xla.DeviceArray): array of query representations with shape (L q b)
d)
        key (jax.interpreters.xla.DeviceArray): array of key representations with shape (L k by d)
        value (jax.interpreters.xla.DeviceArray): array of value representations with shape (L k b)
d) where L \ v = L \ k
       mask (jax.interpreters.xla.DeviceArray): attention-mask, gates attention with shape (L q
by L k)
    Returns:
    jax.interpreters.xla.DeviceArray: Self-attention array for q, k, v arrays. (L_q by L_k)
    assert query.shape[-1] == key.shape[-1] == value.shape[-1], "Embedding dimensions of q, k, v ar
en't all the same"
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # Save depth/dimension of the query embedding for scaling down the dot product
    depth = query.shape[-1]
    # Calculate scaled query key dot product according to formula above
    dots = jnp.matmul(query, jnp.swapaxes(key, 1, 2)) / jnp.sqrt(depth)
    # Apply the mask
    if mask is not None: # The 'None' in this line does not need to be replaced
        dots = jnp.where(mask, dots, jnp.full_like(dots, -1e9))
    # Softmax formula implementation
    # Use trax.fastmath.logsumexp of dots to avoid underflow by division by large numbers
    # Hint: Last axis should be used and keepdims should be True
    \# Note: softmax = e^{(dots - logsumexp(dots))} = E^{dots / sumexp(dots)}
    logsumexp = trax.fastmath.logsumexp(dots, axis=-1, keepdims=True)
    # Take exponential of dots minus logsumexp to get softmax
    # Use jnp.exp()
    dots = jnp.exp(dots - logsumexp)
    # Multiply dots by value to get self-attention
    # Use jnp.matmul()
    attention = jnp.matmul(dots, value)
    ## END CODE HERE ###
    return attention
DotProductAttention(q with batch, k with batch, v with batch, m bool)
Out[58]:
DeviceArray([[[0.8496746 , 0.15032545, 0.8496746 ],
                     , 0.
                                , 1. ]]], dtype=float32)
Expected Output:
```

2.2 Causal Attention

Now you are going to implement causal attention: multi-headed attention with a mask to attend only to words that occurred before.

In the image above, a word can see everything that is before it, but not what is after it. To implement causal attention, you will have to transform vectors and do many reshapes. You will need to implement the functions below.

Exercise 02

Implement the following functions that will be needed for Causal Attention:

- compute_attention_heads: Gets an input \$x\$ of dimension (batch_size, seqlen, n_heads \$\times\$ d_head) and splits the last (depth) dimension and stacks it to the zeroth dimension to allow matrix multiplication (batch_size \$\times\$ n_heads, seqlen, d_head).
- dot_product_self_attention : Creates a mask matrix with False values above the diagonal and True values below and calls DotProductAttention which implements dot product self attention.
- compute_attention_output : Undoes compute_attention_heads by splitting first (vertical) dimension and stacking in the last (depth) dimension (batch_size, seqlen, n_heads \$\times\$ d_head). These operations concatenate (stack/merge) the heads.

Next there are some toy tensors which may serve to give you an idea of the data shapes and opperations involved in Causal Attention. They are also useful to test out your functions!

```
In [59]:
```

```
tensor2d = create tensor(q)
display tensor(tensor2d, 'query matrix (2D tensor)')
tensor4d2b = create tensor([[q, q], [q, q]])
display tensor(tensor4d2b, 'batch of two (multi-head) collections of query matrices (4D tensor)')
tensor3dc = create tensor([jnp.concatenate([q, q], axis = -1)])
display tensor(tensor3dc, 'one batch of concatenated heads of query matrices (3d tensor)')
tensor3dc3b = create tensor([jnp.concatenate([q, q], axis = -1), jnp.concatenate([q, q], axis = -1)
, jnp.concatenate([q, q], axis = -1)])
display tensor(tensor3dc3b, 'three batches of concatenated heads of query matrices (3d tensor)')
query matrix (2D tensor) shape: (2, 3)
[[1 0 0]
 [0 1 0]]
batch of two (multi-head) collections of query matrices (4D tensor) shape: (2, 2, 2, 3)
[[[[1 0 0]
  [0 1 0]]
  [[1 0 0]
   [0 1 0]]
 [[[1 0 0]
   [0 1 0]]
  [[1 0 0]
   [0 1 0]]]
one batch of concatenated heads of query matrices (3d tensor) shape: (1, 2, 6)
[[[1 0 0 1 0 0]
  [0 1 0 0 1 0]]]
three batches of concatenated heads of query matrices (3d tensor) shape: (3, 2, 6)
[[[1 0 0 1 0 0]
 [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
 [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
 [0 1 0 0 1 0]]]
```

It is important to know that the following 3 functions would normally be defined within the CausalAttention function further below.

However this makes these functions harder to test. Because of this, these functions are shown individually using a closure (when necessary) that simulates them being inside of the CausalAttention function. This is done because they rely on some variables

Support Functions

compute_attention_heads: Gets an input \$x\$ of dimension (batch_size, seqlen, n_heads \$\times\$ d_head) and splits the last (depth) dimension and stacks it to the zeroth dimension to allow matrix multiplication (batch_size \$\times\$ n_heads, seqlen, d_head).

For the closures you only have to fill the inner function.

```
In [64]:
```

```
# UNQ C2
# GRADED FUNCTION: compute attention heads closure
def compute attention heads closure (n heads, d head):
    """ Function that simulates environment inside CausalAttention function.
       d head (int): dimensionality of heads.
       n heads (int): number of attention heads.
   Returns:
       function: compute attention heads function
   def compute attention heads(x):
        """ Compute the attention heads.
           x (jax.interpreters.xla.DeviceArray): tensor with shape (batch_size, seqlen, n heads X
d head).
           jax.interpreters.xla.DeviceArray: reshaped tensor with shape (batch_size X n_heads,
seqlen, d_head).
       ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
        # Size of the x's batch dimension
       batch size = x.shape[0]
       # Length of the sequence
        # Should be size of x's first dimension without counting the batch dim
       seqlen = x.shape[1]
       # Reshape x using jnp.reshape()
        # batch_size, seqlen, n_heads*d_head -> batch_size, seqlen, n_heads, d_head
       x = jnp.reshape(x, (batch size, seqlen, n heads, d head))
        # Transpose x using jnp.transpose()
        # batch_size, seqlen, n_heads, d_head -> batch_size, n_heads, seqlen, d_head
        # Note that the values within the tuple are the indexes of the dimensions of x and you mus
t rearrange them
       x = jnp.transpose(x, (0, 2, 1, 3))
        # Reshape x using jnp.reshape()
       # batch size, n heads, seqlen, d head -> batch size*n heads, seqlen, d head
       x = jnp.reshape(x, (batch size*n heads, seqlen, d head))
       ### END CODE HERE ###
       return x
   return compute attention heads
```

```
In [65]:
```

```
display_tensor(tensor3dc3b, "input tensor")
result_cah = compute_attention_heads_closure(2,3)(tensor3dc3b)
display_tensor(result_cah, "output tensor")

input tensor shape: (3, 2, 6)

[[[1 0 0 1 0 0]
       [0 1 0 0 1 0]]

[[1 0 0 1 0 0]
       [0 1 0 0 1 0]]

[[1 0 0 1 0 0]
       [0 1 0 0 1 0]]]

output tensor shape: (6, 2, 3)
```

```
[[[1 0 0]

[0 1 0]]]

[[1 0 0]

[0 1 0]]

[[1 0 0]

[0 1 0]]

[[1 0 0]

[0 1 0]]

[[1 0 0]

[[0 1 0]]

[[1 0 0]

[[0 1 0]]
```

```
input tensor shape: (3, 2, 6)
[[[1 0 0 1 0 0]
  [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
  [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
  [0 1 0 0 1 0]]]
output tensor shape: (6, 2, 3)
[[[1 0 0]
  [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
 [0 1 0]]]
```

dot_product_self_attention : Creates a mask matrix with False values above the diagonal and True values below and calls DotProductAttention which implements dot product self attention.

```
In [68]:
```

```
### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

# Hint: mask size should be equal to L_q. Remember that q has shape (batch_size, L_q, d)
mask_size = q.shape[1]

# Creates a matrix with ones below the diagonal and 0s above. It should have shape (1, mask_si
ze, mask_size)
# Notice that 1's and 0's get casted to True/False by setting dtype to jnp.bool_
# Use jnp.tril() - Lower triangle of an array and jnp.ones()
mask = jnp.tril(jnp.ones((1, mask_size, mask_size), dtype=jnp.bool_), k=0)

### END CODE HERE ###

return DotProductAttention(q, k, v, mask)
```

In [69]:

```
dot_product_self_attention(q_with_batch, k_with_batch, v_with_batch)
```

Out[69]:

Expected Output:

```
DeviceArray([[[0. , 1. , 0. ], [0.8496746 , 0.15032543, 0.8496746 ]]], dtype=float32)
```

compute_attention_output: Undoes compute_attention_heads by splitting first (vertical) dimension and stacking in the last (depth) dimension (batch size, seglen, n heads \$\times\$ d head). These operations concatenate (stack/merge) the heads.

In [80]:

```
# UNQ C4
# GRADED FUNCTION: compute_attention_output_closure
def compute attention output closure (n heads, d head):
   """ Function that simulates environment inside CausalAttention function.
   Aras:
       d head (int): dimensionality of heads.
       n_heads (int): number of attention heads.
   Returns:
       function: compute_attention_output function
   def compute attention output(x):
        """ Compute the attention output.
            x (jax.interpreters.xla.DeviceArray): tensor with shape (batch size X n heads, seqlen,
d head).
           jax.interpreters.xla.DeviceArray: reshaped tensor with shape (batch size, seqlen, n heat
ds \ X \ d \ head).
        ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
       # Length of the sequence
       # Should be size of x's first dimension without counting the batch dim
       seqlen = x.shape[1]
       # Reshape x using jnp.reshape() to shape (batch_size, n_heads, seqlen, d_head)
       x = jnp.reshape(x, (-1, n heads, seqlen, d head))
        # Transpose x using jnp.transpose() to shape (batch size, seqlen, n heads, d head)
       x = jnp.transpose(x, (0,2,1,3))
        ### END CODE HERE ###
        # Reshape to allow to concatenate the heads
       return jnp.reshape(x, (-1, seqlen, n_heads * d_head))
   return compute attention output
```

```
In [81]:
display_tensor(result_cah, "input tensor")
result_cao = compute_attention_output_closure(2,3)(result_cah)
display_tensor(result_cao, "output tensor")
input tensor shape: (6, 2, 3)
[[[1 0 0]
 [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
 [0 1 0]]]
output tensor shape: (3, 2, 6)
[[[1 0 0 1 0 0]
 [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
 [0 1 0 0 1 0]]
```

[[1 0 0 1 0 0] [0 1 0 0 1 0]]]

```
input tensor shape: (6, 2, 3)
[[[1 0 0]
  [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
  [0 1 0]]
 [[1 0 0]
 [0 1 0]]
 [[1 0 0]
  [0 1 0]]]
output tensor shape: (3, 2, 6)
[[[1 0 0 1 0 0]
 [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
  [0 1 0 0 1 0]]
 [[1 0 0 1 0 0]
```

Causal Attention Function

Now it is time for you to put everything together within the CausalAttention or Masked multi-head attention function:

Instructions: Implement the causal attention. Your model returns the causal attention through a \$tl.Serial\$ with the following:

- [tl.Branch](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Branch): consisting of 3 [tl.Dense(d feature), ComputeAttentionHeads] to account for the queries, keys, and values.
- [tl.Fn](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.base.Fn): Takes in dot_product_self_attention function and uses it to compute the dot product using \$Q\$, \$K\$, \$V\$.
- [tl.Fn](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.base.Fn): Takes in compute_attention_output_closure to allow for parallel computing.
- [tl.Dense](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dense): Final Dense layer, with dimension d feature.

Remember that in order for trax to properly handle the functions you just defined, they need to be added as layers using the tl.Fn() function.

In [84]:

```
# UNQ C5
# GRADED FUNCTION: CausalAttention
def CausalAttention (d feature,
                    n heads,
                    compute attention heads closure-compute attention heads closure,
                    dot_product_self_attention=dot_product_self_attention,
                    compute attention output closure=compute attention output closure,
                    mode='train'):
    """Transformer-style multi-headed causal attention.
    Aras:
       d feature (int): dimensionality of feature embedding.
       n heads (int): number of attention heads.
       compute attention heads closure (function): Closure around compute attention heads.
       dot product self attention (function): dot product self attention function.
       compute attention output closure (function): Closure around compute attention output.
       mode (str): 'train' or 'eval'.
    Returns:
       trax.layers.combinators.Serial: Multi-headed self-attention model.
    assert d feature % n heads == 0
    d_head = d_feature // n_heads
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # HINT: The second argument to tl.Fn() is an uncalled function (without the parentheses)
    # Since you are dealing with closures you might need to call the outer
    # function with the correct parameters to get the actual uncalled function.
    ComputeAttentionHeads = tl.Fn('AttnHeads', compute_attention_heads_closure(n_heads, d_head), n_
out=1)
    return tl.Serial(
       tl.Branch( # creates three towers for one input, takes activations and creates queries
kevs and values
            [tl.Dense(d_feature), ComputeAttentionHeads], # queries
            [tl.Dense(d_feature), ComputeAttentionHeads], \# keys
            [tl.Dense(d feature), ComputeAttentionHeads], # values
       ),
        tl.Fn('DotProductAttn', dot product self attention, n out=1), # takes QKV
        # HINT: The second argument to tl.Fn() is an uncalled function
        # Since you are dealing with closures you might need to call the outer
        # function with the correct parameters to get the actual uncalled function.
       tl.Fn('AttnOutput', compute_attention_output_closure(n_heads, d_head), n_out=1), # to allow
for parallel
       tl.Dense(d_feature) # Final dense layer
```

```
### END CODE HERE ###

In [85]:

# Take a look at the causal attention model
print(CausalAttention(d_feature=512, n_heads=8))

Serial[
Branch_out3[
[Dense_512, AttnHeads]
[Dense_512, AttnHeads]
[Dense_512, AttnHeads]
```

AttnOutput Dense_512

DotProductAttn in3

```
Serial[

Branch_out3[

[Dense_512, AttnHeads]

[Dense_512, AttnHeads]

[Dense_512, AttnHeads]

]

DotProductAttn_in3

AttnOutput

Dense_512

]
```

2.3 Transformer decoder block

Now that you have implemented the causal part of the transformer, you will implement the transformer decoder block. Concretely you will be implementing this image now.

To implement this function, you will have to call the CausalAttention or Masked multi-head attention function you implemented above. You will have to add a feedforward which consists of:

- [tl.LayerNorm](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.normalization.LayerNorm): used to layer normalize
- [tl.Dense](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dense): the dense layer
- [ff_activation](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.activation_fns.Relu): feed forward activation (we use ReLu) here.
- [tl.Dropout](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dropout): dropout layer
- [tl.Dense](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dense): dense layer
- [tl.Dropout](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dropout): dropout layer

Finally once you implement the feedforward, you can go ahead and implement the entire block using:

- [tl.Residual](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Residual): takes in the tl.LayerNorm(), causal attention block, tl.dropout.
- [tl.Residual](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Residual): takes in the feedforward block you will implement.

Exercise 03

Instructions: Implement the transformer decoder block. Good luck!

In [88]:

```
# UNQ C6
# GRADED FUNCTION: DecoderBlock
def DecoderBlock(d_model, d_ff, n_heads,
                 dropout, mode, ff_activation):
    """Returns a list of layers that implements a Transformer decoder block.
    The input is an activation tensor.
    Aras:
        d model (int): depth of embedding.
        d ff (int): depth of feed-forward layer.
        n heads (int): number of attention heads.
        dropout (float): dropout rate (how much to drop out).
       mode (str): 'train' or 'eval'.
        ff activation (function): the non-linearity in feed-forward layer.
    Returns:
       list: list of trax.layers.combinators.Serial that maps an activation tensor to an
activation tensor.
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # Create masked multi-head attention block using CausalAttention function
    causal attention = CausalAttention(
                        d model,
                        n heads=n heads,
                        mode=mode
    # Create feed-forward block (list) with two dense layers with dropout and input normalized
    feed forward = [
        # Normalize layer inputs
        tl.LayerNorm(),
        # Add first feed forward (dense) layer (don't forget to set the correct value for n units)
        tl.Dense(d ff),
        # Add activation function passed in as a parameter (you need to call it!)
       ff activation(), # Generally ReLU
        # Add dropout with rate and mode specified (i.e., don't use dropout during evaluation)
       tl.Dropout (rate=dropout, mode=mode),
        # Add second feed forward layer (don't forget to set the correct value for n_units)
       tl.Dense(d model),
         Add dropout with rate and mode specified (i.e., don't use dropout during evaluation)
        tl.Dropout(rate=dropout, mode=mode),
    ]
    # Add list of two Residual blocks: the attention with normalization and dropout and feed-forwa
rd blocks
    return [
      tl.Residual(
         # Normalize layer input
         tl.LayerNorm(),
          # Add causal attention block previously defined (without parentheses)
         causal attention,
          # Add dropout with rate and mode specified
         tl.Dropout(rate=dropout, mode=mode)
       ),
      tl.Residual(
          # Add feed forward block (without parentheses)
          feed forward
        ),
```

```
### END CODE HERE ###
```

In [89]:

```
# Take a look at the decoder block
print(DecoderBlock(d model=512, d ff=2048, n heads=8, dropout=0.1, mode='train', ff activation=tl.R
elu))
[Serial[
 Branch_out2[
   None
   Serial[
     LayerNorm
     Serial[
        Branch out3[
          [Dense_512, AttnHeads]
          [Dense_512, AttnHeads]
          [Dense 512, AttnHeads]
        DotProductAttn in3
       AttnOutput
       Dense_512
     Dropout
    ]
 Add_in2
], Serial[
 Branch_out2[
   None
   Serial[
     LayerNorm
     Dense 2048
     Relu
     Dropout
     Dense 512
     Dropout
   ]
 Add_in2
]]
```

Expected Output:

```
[Serial[
  Branch_out2[
   None
    Serial[
      LayerNorm
      Serial[
        Branch out3[
          [Dense_512, AttnHeads]
          [Dense_512, AttnHeads]
          [Dense_512, AttnHeads]
        ]
        DotProductAttn_in3
        AttnOutput
        Dense_512
      ]
      Dropout
    ]
  Add in2
], Serial[
  Branch_out2[
    None
    Serial[
      LayerNorm
      Dense 2048
```

```
Relu
Dropout
Dense_512
Dropout

]
Add_in2
]
```

2.4 Transformer Language Model

You will now bring it all together. In this part you will use all the subcomponents you previously built to make the final model. Concretely, here is the image you will be implementing.

Exercise 04

Instructions: Previously you coded the decoder block. Now you will code the transformer language model. Here is what you will need.

- positional_enconder a list containing the following layers:
 - tl.Embedding
 - tl.Dropout
 - tl.PositionalEncoding
- A list of n_layers decoder blocks.
- [tl.Serial](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Serial): takes in the following layers or lists of layers:
 - [tl.ShiftRight](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.attention.ShiftRight): : shift the tensor to the right by padding on axis 1.
 - positional encoder: encodes the text positions.
 - decoder_blocks : the ones you created.
 - $\qquad \textbf{[tl.LayerNorm]} (\texttt{https://trax-ml.readthedocs.io/en/latest/trax.layers.html\#trax.layers.normalization.LayerNorm): a layer norm.} \\$
 - [tl.Dense](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Dense): takes in the vocab_size.
 - [tl.LogSoftmax](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.LogSoftmax): to predict.

Go go go!! You can do it :)

```
In [92]:
```

```
aropout=0.1,
                  max len=4096,
                  mode='train',
                 ff activation=tl.Relu):
    """Returns a Transformer language model.
    The input to the model is a tensor of tokens. (This model uses only the
    decoder part of the overall Transformer.)
    Args:
        vocab size (int): vocab size.
        d model (int): depth of embedding.
        d ff (int): depth of feed-forward layer.
        n layers (int): number of decoder layers.
        n heads (int): number of attention heads.
        dropout (float): dropout rate (how much to drop out).
       max len (int): maximum symbol length for positional encoding.
        mode (str): 'train', 'eval' or 'predict', predict mode is for fast inference.
        ff_activation (function): the non-linearity in feed-forward layer.
    Returns:
       trax.layers.combinators.Serial: A Transformer language model as a layer that maps from a t
ensor of tokens
       to activations over a vocab set.
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # Embedding inputs and positional encoder
    positional encoder = [
        # Add embedding layer of dimension (vocab size, d model)
        tl.Embedding(vocab size, d model),
        # Use dropout with rate and mode specified
        tl.Dropout(rate=dropout, mode=mode),
        # Add positional encoding layer with maximum input length and mode specified
        tl.PositionalEncoding(max len=max len, mode=mode)]
    # Create stack (list) of decoder blocks with n layers with necessary parameters
    decoder blocks = [
        DecoderBlock(d model, d ff, n heads, dropout, mode, ff activation) for in range(n layers)
    # Create the complete model as written in the figure
    return tl.Serial(
        # Use teacher forcing (feed output of previous step to current step)
        tl.ShiftRight(mode=mode), # Specify the mode!
        # Add positional encoder
       positional encoder,
        # Add decoder blocks
       decoder_blocks,
        # Normalize layer
        tl.LayerNorm(),
        # Add dense layer of vocab size (since need to select a word to translate to)
        # (a.k.a., logits layer. Note: activation already set by ff activation)
        tl.Dense(vocab_size),
        # Get probabilities with Logsoftmax
        tl.LogSoftmax()
    ### END CODE HERE ###
In [93]:
```

```
# Take a look at the Transformer
print(TransformerLM(n_layers=1))

Serial[
    ShiftRight(1)
    Embedding_33300_512
    Dropout
    PositionalEncoding
    Serial[
         Branch_out2[
         None
         Serial[
```

```
LayerNorm
         Serial[
          Branch_out3[
             [Dense_512, AttnHeads]
[Dense_512, AttnHeads]
             [Dense_512, AttnHeads]
          DotProductAttn_in3
           AttnOutput
          Dense_512
         Dropout
      ]
    Add_in2
  ]
  Serial[
    Branch_out2[
      None
      Serial[
        LayerNorm
        Dense_2048
        Relu
        Dropout
        Dense 512
         Dropout
      ]
    Add_in2
  ]
  LayerNorm
  Dense_33300
  LogSoftmax
]
```

```
Serial[
    ShiftRight(1)
    Embedding_33300_512
    Dropout
    PositionalEncoding
     Serial[
       Branch_out2[
        None
         Serial[
          LayerNorm
          Serial[
             Branch_out3[
               [Dense_512, AttnHeads]
               [Dense_512, AttnHeads]
               [Dense_512, AttnHeads]
             DotProductAttn_in3
             AttnOutput
             Dense_512
          ]
          Dropout
         ]
       ]
       Add_in2
     ]
     Serial[
       Branch out2[
        None
        Serial[
          LayerNorm
          Dense_2048
          Relu
          Dropout
```

```
Dense_512
Dense_512
Dropout

|
|
|
|
| Add_in2
|
|
| LayerNorm
| Dense_33300
| LogSoftmax
```

Part 3: Training

Now you are going to train your model. As usual, you have to define the cost function, the optimizer, and decide whether you will be training it on a <code>gpu</code> or <code>cpu</code>. In this case, you will train your model on a cpu for a few steps and we will load in a pre-trained model that you can use to predict with your own words.

3.1 Training the model

You will now write a function that takes in your model and trains it. To train your model you have to decide how many times you want to iterate over the entire data set. Each iteration is defined as an epoch. For each epoch, you have to go over all the data, using your training iterator.

Exercise 05

Instructions: Implement the train model program below to train the neural network above. Here is a list of things you should do:

- Create the train task by calling <u>trax.supervised.training.TrainTask</u> and pass in the following:
 - labeled_data = train_gen
 - loss_fn = <u>tl.CrossEntropyLoss()</u>
 - optimizer = trax.optimizers.Adam(0.01)
 - Ir_schedule = Ir_schedule
- Create the eval task by calling trax.supervised.training.EvalTask and pass in the following:
 - labeled data = eval gen
 - metrics = tl.CrossEntropyLoss() and <u>tl.Accuracy()</u>
- Create the training loop by calling trax.supervised.Training.Loop and pass in the following:
 - TransformerLM
 - train_task
 - eval_task = [eval_task]
 - output_dir = output_dir

You will be using a cross entropy loss, with Adam optimizer. Please read the <u>Trax</u> documentation to get a full understanding.

The training loop that this function returns can be runned using the run () method by passing in the desired number of steps.

In [94]:

```
lr_schedule = trax.lr.warmup_and_rsqrt_decay(n_warmup_steps=1000, max_value=0.01)
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
   train task = training.TrainTask(
     labeled data=train gen, # The training generator
     loss_layer=tl.CrossEntropyLoss(), # Loss function
     optimizer=trax.optimizers.Adam(0.01), # Optimizer (Don't forget to set LR to 0.01)
     lr schedule=lr schedule,
     n steps per checkpoint=10
   eval_task = training.EvalTask(
     labeled data=eval gen, # The evaluation generator
     metrics=[tl.CrossEntropyLoss(), tl.CrossEntropyLoss() and tl.Accuracy()] # CrossEntropyLoss a
nd Accuracy
   ### END CODE HERE ###
   loop = training.Loop(TransformerLM(d model=4,
                                       n layers=1,
                                       n_heads=2,
                                       mode='train'),
                         train task,
                         eval tasks=[eval task],
                         output dir=output dir)
   return loop
```

Notice that the model will be trained for only 10 steps.

Even with this constraint the model with the original default arguments took a very long time to finish. Because of this some parameters are changed when defining the model that is fed into the training loop in the function above.

In [95]:

```
# Should take around 1.5 minutes
!rm -f ~/model/model.pkl.gz
loop = training loop(TransformerLM, train batch stream, eval batch stream)
loop.run(10)
           1: Ran 1 train steps in 11.09 secs
Step
           1: train CrossEntropyLoss | 10.41245461
1: eval CrossEntropyLoss | 10.41230392
Step
Step
                               Accuracy | 0.00000000
          1: eval
Step
          10: Ran 9 train steps in 55.12 secs
Step
          10: train CrossEntropyLoss | 10.41435528
10: eval CrossEntropyLoss | 10.41224480
Step
Step
                               Accuracy | 0.00000000
Step
          10: eval
```

Part 4: Evaluation

4.1 Loading in a trained model

In this part you will evaluate by loading in an almost exact version of the model you coded, but we trained it for you to save you time. Please run the cell below to load in the model.

As you may have already noticed the model that you trained and the pretrained model share the same overall architecture but they have different values for some of the parameters:

```
Original (pretrained) model:

TransformerLM(vocab_size=33300, d_model=512, d_ff=2048, n_layers=6, n_heads=8, dropout=0.1, max len=4096, ff activation=tl.Relu)
```

```
Your model:
```

```
TransformerLM(d model=4, d ff=16, n layers=1, n heads=2)
```

Only the parameters shown for your model were changed. The others stayed the same.

```
In [96]:
```

```
# Get the model architecture
model = TransformerLM(mode='eval')

# Load the pre-trained weights
model.init_from_file('model.pkl.gz', weights_only=True)
```

Part 5: Testing with your own input

You will now test your input. You are going to implement greedy decoding. This consists of two functions. The first one allows you to identify the next symbol. It gets the argmax of the output of your model and then returns that index.

Exercise 06

Instructions: Implement the next symbol function that takes in the cur_output_tokens and the trained model to return the index of the

```
In [97]:
```

```
# UNO C9
def next symbol(cur output tokens, model):
    """Returns the next symbol for a given sentence.
   Args:
       cur output tokens (list): tokenized sentence with EOS and PAD tokens at the end.
       model (trax.layers.combinators.Serial): The transformer model.
   Returns:
       int: tokenized symbol.
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
   # current output tokens length
   token length = len(cur output_tokens)
    # calculate the minimum power of 2 big enough to store token length
    # HINT: use np.ceil() and np.log2()
    # add 1 to token length so np.log2() doesn't receive 0 when token length is 0
   padded_length = 2**int(np.ceil(np.log2(token_length + 1)))
    # Fill cur output tokens with 0's until it reaches padded length
   padded = cur_output_tokens + [0] * (padded_length - token_length)
   padded with batch = np.array(padded)[None, :] # Don't replace this 'None'! This is a way of set
ting the batch dim
    # model expects a tuple containing two padded tensors (with batch)
   output, _ = model((padded_with_batch, padded with batch))
    # HINT: output has shape (1, padded length, vocab size)
    # To get log probs you need to index output with 0 in the first dim
    # token length in the second dim and all of the entries for the last dim.
   log probs = output[0, token length, :]
    ### END CODE HERE ###
   return int(np.argmax(log probs))
```

```
In [98]:
```

```
# Test it out!
sentence_test_nxt_symbl = "I want to fly in the sky."
detokenize([next_symbol(tokenize(sentence_test_nxt_symbl)+[0], model)])
```

```
'The'
```

```
'The'
```

5.1 Greedy decoding

Now you will implement the greedy_decode algorithm that will call the next_symbol function. It takes in the input_sentence, the trained model and returns the decoded sentence.

Exercise 07

Instructions: Implement the greedy_decode algorithm.

In [103]:

```
# UNQ C10
# Decoding functions.
def greedy decode(input sentence, model):
    """Greedy decode function.
       input sentence (string): a sentence or article.
       model (trax.layers.combinators.Serial): Transformer model.
    Returns:
       string: summary of the input.
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # Use tokenize()
    cur_output_tokens = tokenize(input sentence) + [0]
    generated output = []
    cur_output = 0
    EOS = 1
    while cur_output != EOS:
       # Get next symbol
       cur output = next symbol(cur output tokens, model)
       # Append next symbol to original sentence
       cur output tokens.append(cur output)
        # Append next symbol to generated sentence
        generated output.append(cur output)
        print(detokenize(generated_output))
    ### END CODE HERE ###
    return detokenize(generated output)
```

In [104]:

: I
: I just
: I just found
: I just found ros
: I just found roses
: I just found roses.

```
# Test it out on a sentence!
test_sentence = "It was a sunny day when I went to the market to buy some flowers. But I only foun
d roses, not tulips."
print(wrapper.fill(test_sentence), '\n')
print(greedy_decode(test_sentence, model))

It was a sunny day when I went to the market to buy some flowers. But
I only found roses, not tulips.
:
```

```
: I just found roses, not
: I just found roses, not tu
: I just found roses, not tulips
: I just found roses, not tulips
: I just found roses, not tulips
: I just found roses, not tulips.
```

```
:
: I
: I just
: I just found
: I just found ros
: I just found roses
: I just found roses,
: I just found roses,
: I just found roses, not
: I just found roses, not tu
: I just found roses, not tulips
: I just found roses, not tulips.
```

In [105]:

```
# Test it out with a whole article!
article = "It's the posing craze sweeping the U.S. after being brought to fame by skier Lindsey Vo
nn, soccer star Omar Cummings, baseball player Albert Pujols - and even Republican politician Rick
Perry. But now four students at Riverhead High School on Long Island, New York, have been suspende
d for dropping to a knee and taking up a prayer pose to mimic Denver Broncos quarterback Tim Tebow
. Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were all suspended for one day be
cause the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to students. Sc
roll down for video. Banned: Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll (all p
ictured left) were all suspended for one day by Riverhead High School on Long Island, New York, fo
r their tribute to Broncos quarterback Tim Tebow. Issue: Four of the pupils were suspended for one
day because they allegedly did not heed to warnings that the 'Tebowing' craze at the school was bl
ocking the hallway and presenting a safety hazard to students."
print(wrapper.fill(article), '\n')
print(greedy_decode(article, model))
```

It's the posing craze sweeping the U.S. after being brought to fame by skier Lindsey Vonn, soccer star Omar Cummings, baseball player Albert Pujols - and even Republican politician Rick Perry. But now four students at Riverhead High School on Long Island, New York, have been suspended for dropping to a knee and taking up a prayer pose to mimic Denver Broncos quarterback Tim Tebow. Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were all suspended for one day because the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to students. Scroll down for video. Banned: Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll (all pictured left) were all suspended for one day by Riverhead High School on Long Island, New York, for their tribute to Broncos quarterback Tim Tebow. Issue: Four of the pupils were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze at the school was blocking the hallway and presenting a safety hazard to students.

```
Jordan
Jordan Ful
Jordan Fulcol
Jordan Fulcoly
Jordan Fulcoly,
Jordan Fulcoly, Wayne
Jordan Fulcoly, Wayne Dre
Jordan Fulcoly, Wayne Drexe
Jordan Fulcoly, Wayne Drexel
Jordan Fulcoly, Wayne Drexel,
Jordan Fulcoly, Wayne Drexel, Tyler
Jordan Fulcoly, Wayne Drexel, Tyler Carroll
```

```
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day.
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not hee
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warn
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that the
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that the '
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that the 'Te
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that the 'Tebow
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended for one day. Four students were suspended for one day
because they allegedly did not heed to warnings that the 'Tebowing
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
```

suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' cra

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocki

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hall

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to students

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to students.

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to students.<

Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were suspended for one day. Four students were suspended for one day because they allegedly did not heed to warnings that the 'Tebowing' craze was blocking the hallway and presenting a safety hazard to

```
Jordan
Jordan Ful
Jordan Fulcol
Jordan Fulcoly
Jordan Fulcoly,
Jordan Fulcoly, Wayne
Jordan Fulcoly, Wayne Dre
Jordan Fulcoly, Wayne Drexe
Jordan Fulcoly, Wayne Drexel
Jordan Fulcoly, Wayne Drexel,
Final summary:
Jordan Fulcoly, Wayne Drexel, Tyler Carroll and Connor Carroll were
suspended {\bf for} one day. Four students were suspended {\bf for} one day
because they allegedly did not heed to warnings that the 'Tebowing'
craze was blocking the hallway and presenting a safety hazard to
students.<EOS>
```

Congratulations on finishing this week's assignment! You did a lot of work and now you should have a better understanding of the encoder part of Transformers and how Transformers can be used for text summarization.

Keep it up!