



Now that we've seen the concepts behind **transfer learning**, let's dig in and take a look at how to do it for ourselves with TensorFlow and Keras.

In the next few videos you'll be using this notebook to explore transfer learning:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/Course%20-%20Part%206%20-%20Lesson%203%20-%20Notebook.ipynb>

For more on how to freeze/lock layers, explore the documentation, which includes an example using MobileNet architecture: https://www.tensorflow.org/tutorials/images/transfer_learning

```
import os

from tensorflow.keras import layers
from tensorflow.keras import Model
```

Transfer Learning

```
https://storage.googleapis.com/mledu-datasets/
inception_v3_weights_tf_dim_ordering_tf_kernels
```

A copy of the pretrained weights for the inception neural network is saved at this URL. Think of this as a snapshot of the model after being trained. It's the parameters that can then get loaded into the skeleton of the model, to turn it back into a trained model.

```
from tensorflow.keras.applications.inception_v3 import InceptionV3

local_weights_file = '/tmp/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5'

pre_trained_model = InceptionV3(input_shape = (150, 150, 3),
                                include_top = False,
                                weights = None)

pre_trained_model.load_weights(local_weights_file)
```

So now if we want to use inception, it's fortunate that keras has the model definition built in. So you instantiate that with the desired input shape for your data, and specify that you don't want to use the built-in weights, but the snapshot that you've just downloaded. The inception V3 has a fully-connected layer at the top. So by setting include_top to false, you're specifying that you want to ignore this and get straight to the convolutions.

```
for layer in pre_trained_model.layers:
    layer.trainable = False
```

Now that I have my pretrained model instantiated, I can iterate through its layers and lock them, saying that they're not going to be trainable with this code.

```
pre_trained_model.summary()
```

InceptionV3				
layer	type	shape	size	name

input	Input	(1, 299, 299, 3)	26544	input
conv2d_1	Conv2D	(3, 3, 32, 3)	864	conv2d_1
conv2d_2	Conv2D	(3, 3, 32, 3)	864	conv2d_2
conv2d_3	Conv2D	(3, 3, 32, 3)	864	conv2d_3
conv2d_4	Conv2D	(3, 3, 32, 3)	864	conv2d_4
conv2d_5	Conv2D	(3, 3, 32, 3)	864	conv2d_5
conv2d_6	Conv2D	(3, 3, 32, 3)	864	conv2d_6
conv2d_7	Conv2D	(3, 3, 32, 3)	864	conv2d_7
conv2d_8	Conv2D	(3, 3, 32, 3)	864	conv2d_8
conv2d_9	Conv2D	(3, 3, 32, 3)	864	conv2d_9
conv2d_10	Conv2D	(3, 3, 32, 3)	864	conv2d_10
conv2d_11	Conv2D	(3, 3, 32, 3)	864	conv2d_11
conv2d_12	Conv2D	(3, 3, 32, 3)	864	conv2d_12
conv2d_13	Conv2D	(3, 3, 32, 3)	864	conv2d_13
conv2d_14	Conv2D	(3, 3, 32, 3)	864	conv2d_14
conv2d_15	Conv2D	(3, 3, 32, 3)	864	conv2d_15
conv2d_16	Conv2D	(3, 3, 32, 3)	864	conv2d_16
conv2d_17	Conv2D	(3, 3, 32, 3)	864	conv2d_17
conv2d_18	Conv2D	(3, 3, 32, 3)	864	conv2d_18
conv2d_19	Conv2D	(3, 3, 32, 3)	864	conv2d_19
conv2d_20	Conv2D	(3, 3, 32, 3)	864	conv2d_20
conv2d_21	Conv2D	(3, 3, 32, 3)	864	conv2d_21
conv2d_22	Conv2D	(3, 3, 32, 3)	864	conv2d_22
conv2d_23	Conv2D	(3, 3, 32, 3)	864	conv2d_23
conv2d_24	Conv2D	(3, 3, 32, 3)	864	conv2d_24
conv2d_25	Conv2D	(3, 3, 32, 3)	864	conv2d_25
conv2d_26	Conv2D	(3, 3, 32, 3)	864	conv2d_26
conv2d_27	Conv2D	(3, 3, 32, 3)	864	conv2d_27
conv2d_28	Conv2D	(3, 3, 32, 3)	864	conv2d_28
conv2d_29	Conv2D	(3, 3, 32, 3)	864	conv2d_29
conv2d_30	Conv2D	(3, 3, 32, 3)	864	conv2d_30
conv2d_31	Conv2D	(3, 3, 32, 3)	864	conv2d_31
conv2d_32	Conv2D	(3, 3, 32, 3)	864	conv2d_32
conv2d_33	Conv2D	(3, 3, 32, 3)	864	conv2d_33
conv2d_34	Conv2D	(3, 3, 32, 3)	864	conv2d_34
conv2d_35	Conv2D	(3, 3, 32, 3)	864	conv2d_35
conv2d_36	Conv2D	(3, 3, 32, 3)	864	conv2d_36
conv2d_37	Conv2D	(3, 3, 32, 3)	864	conv2d_37
conv2d_38	Conv2D	(3, 3, 32, 3)	864	conv2d_38
conv2d_39	Conv2D	(3, 3, 32, 3)	864	conv2d_39
conv2d_40	Conv2D	(3, 3, 32, 3)	864	conv2d_40
conv2d_41	Conv2D	(3, 3, 32, 3)	864	conv2d_41
conv2d_42	Conv2D	(3, 3, 32, 3)	864	conv2d_42
conv2d_43	Conv2D	(3, 3, 32, 3)	864	conv2d_43
conv2d_44	Conv2D	(3, 3, 32, 3)	864	conv2d_44
conv2d_45	Conv2D	(3, 3, 32, 3)	864	conv2d_45
conv2d_46	Conv2D	(3, 3, 32, 3)	864	conv2d_46
conv2d_47	Conv2D	(3, 3, 32, 3)	864	conv2d_47
conv2d_48	Conv2D	(3, 3, 32, 3)	864	conv2d_48
conv2d_49	Conv2D	(3, 3, 32, 3)	864	conv2d_49
conv2d_50	Conv2D	(3, 3, 32, 3)	864	conv2d_50
conv2d_51	Conv2D	(3, 3, 32, 3)	864	conv2d_51
conv2d_52	Conv2D	(3, 3, 32, 3)	864	conv2d_52
conv2d_53	Conv2D	(3, 3, 32, 3)	864	conv2d_53
conv2d_54	Conv2D	(3, 3, 32, 3)	864	conv2d_54
conv2d_55	Conv2D	(3, 3, 32, 3)	864	conv2d_55
conv2d_56	Conv2D	(3, 3, 32, 3)	864	conv2d_56
conv2d_57	Conv2D	(3, 3, 32, 3)	864	conv2d_57
conv2d_58	Conv2D	(3, 3, 32, 3)	864	conv2d_58
conv2d_59	Conv2D	(3, 3, 32, 3)	864	conv2d_59
conv2d_60	Conv2D	(3, 3, 32, 3)	864	conv2d_60
conv2d_61	Conv2D	(3, 3, 32, 3)	864	conv2d_61
conv2d_62	Conv2D	(3, 3, 32, 3)	864	conv2d_62
conv2d_63	Conv2D	(3, 3, 32, 3)	864	conv2d_63
conv2d_64	Conv2D	(3, 3, 32, 3)	864	conv2d_64
conv2d_65	Conv2D	(3, 3, 32, 3)	864	conv2d_65
conv2d_66	Conv2D	(3, 3, 32, 3)	864	conv2d_66
conv2d_67	Conv2D	(3, 3, 32, 3)	864	conv2d_67
conv2d_68	Conv2D	(3, 3, 32, 3)	864	conv2d_68
conv2d_69	Conv2D	(3, 3, 32, 3)	864	conv2d_69
conv2d_70	Conv2D	(3, 3, 32, 3)	864	conv2d_70
conv2d_71	Conv2D	(3, 3, 32, 3)	864	conv2d_71
conv2d_72	Conv2D	(3, 3, 32, 3)	864	conv2d_72
conv2d_73	Conv2D	(3, 3, 32, 3)	864	conv2d_73
conv2d_74	Conv2D	(3, 3, 32, 3)	864	conv2d_74
conv2d_75	Conv2D	(3, 3, 32, 3)	864	conv2d_75
conv2d_76	Conv2D	(3, 3, 32, 3)	864	conv2d_76
conv2d_77	Conv2D	(3, 3, 32, 3)	864	conv2d_77
conv2d_78	Conv2D	(3, 3, 32, 3)	864	conv2d_78
conv2d_79	Conv2D	(3, 3, 32, 3)	864	conv2d_79
conv2d_80	Conv2D	(3, 3, 32, 3)	864	conv2d_80
conv2d_81	Conv2D	(3, 3, 32, 3)	864	conv2d_81
conv2d_82	Conv2D	(3, 3, 32, 3)	864	conv2d_82
conv2d_83	Conv2D	(3, 3, 32, 3)	864	conv2d_83
conv2d_84	Conv2D	(3, 3, 32, 3)	864	conv2d_84
conv2d_85	Conv2D	(3, 3, 32, 3)	864	conv2d_85
conv2d_86	Conv2D	(3, 3, 32, 3)	864	conv2d_86
conv2d_87	Conv2D	(3, 3, 32, 3)	864	conv2d_87
conv2d_88	Conv2D	(3, 3, 32, 3)	864	conv2d_88
conv2d_89	Conv2D	(3, 3, 32, 3)	864	conv2d_89
conv2d_90	Conv2D	(3, 3, 32, 3)	864	conv2d_90
conv2d_91	Conv2D	(3, 3, 32, 3)	864	conv2d_91
conv2d_92	Conv2D	(3, 3, 32, 3)	864	conv2d_92
conv2d_93	Conv2D	(3, 3, 32, 3)	864	conv2d_93
conv2d_94	Conv2D	(3, 3, 32, 3)	864	conv2d_94
conv2d_95	Conv2D	(3, 3, 32, 3)	864	conv2d_95
conv2d_96	Conv2D	(3, 3, 32, 3)	864	conv2d_96
conv2d_97	Conv2D	(3, 3, 32, 3)	864	conv2d_97
conv2d_98	Conv2D	(3, 3, 32, 3)	864	conv2d_98
conv2d_99	Conv2D	(3, 3, 32, 3)	864	conv2d_99
conv2d_100	Conv2D	(3, 3, 32, 3)	864	conv2d_100

Adding your DNN

In the previous video you saw how to take the layers from an existing model, and make them so that they don't get retrained -- i.e. you freeze (or lock) the already learned convolutions into your model. Now, you'll need to add your own DNN at the bottom of these, which you can retrain to your data. In the next video you'll see how to do that...

```
last_layer = pre_trained_model.get_layer('mixed7')
```

```
last_output = last_layer.output
```

All of the layers have names, so you can look up the name of the last layer that you want to use. If you inspect the summary, you'll see that the bottom layers have convolved to 3 by 3. But I want to use something with a little more information. So I moved up the model description to find mixed7, which is the output of a lot of convolution that are 7 by 7.

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dense(1, activation='sigmoid')(x)
```

```
model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])
```

we'll define our new model, taking the output from the inception model's mixed7 layer. You start by flattening the input, which just happens to be the output from inception

```

from tensorflow.keras.optimizers import RMSprop

x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dense(1, activation='sigmoid')(x)

model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])

```

```

from tensorflow.keras.optimizers import RMSprop

x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dense(1, activation='sigmoid')(x)

model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])

```

You can then create a model using
 the Model abstract class. And
 passing at the input and the layers
 definition that you've just created.

```

from tensorflow.keras.optimizers import RMSprop

x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dense(1, activation='sigmoid')(x)


model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])

```

```
# Add our data-augmentation parameters to ImageDataGenerator
train_datagen = ImageDataGenerator(rescale = 1./255.,
                                    rotation_range = 40,
                                    width_shift_range = 0.2,
                                    height_shift_range = 0.2,
                                    shear_range = 0.2,
                                    zoom_range = 0.2,
                                    horizontal_flip = True)
```

```
train_generator = train_datagen.flow_from_directory(
    train_dir,
    batch_size = 20,
    class_mode = 'binary',
    target_size = (150, 150))
```

```
history = model.fit_generator(
    train_generator,
    validation_data = validation_generator,
    steps_per_epoch = 100,
    epochs = 100,
    validation_steps = 50,
    verbose = 2)
```






Using dropouts!

Another useful tool to explore at this point is the Dropout.

The idea behind Dropouts is that they remove a random number of neurons in your neural network. This works very well for two reasons: The first is that neighboring neurons often end up with similar weights, which can lead to overfitting, so dropping some out at random can remove this. The second is that often a neuron can over-weigh the input from a neuron in the previous layer, and can over specialize as a result. Thus, dropping out can break the neural network out of this potential bad habit!

Check out Andrew's terrific video explaining dropouts here: <https://www.youtube.com/watch?v=ARq74QuavAo>





What's interesting if you do this, is that you end up with another but a different overfitting situation. Here is the graph of the accuracy of training versus validation. As you can see, while it started out well, the validation is diverging away from the training in a really bad way. So, how do we fix this?

```
from tensorflow.keras.optimizers import RMSprop

x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dense(1, activation='sigmoid')(x)


model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])
```

```
from tensorflow.keras.optimizers import RMSprop

x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.2)(x)
x = layers.Dense(1, activation='sigmoid')(x)

model = Model(pre_trained_model.input, x)
model.compile(optimizer = RMSprop(lr=0.0001),
              loss = 'binary_crossentropy',
              metrics = ['acc'])
```

The parameter is between 0 and 1 and it's the fraction of units to drop. In this case, we're dropping out 20% of our neurons.



Dropout

