




Fashion MNIST


- 70k Images
- 10 Categories
- Images are 28x28
- Can train a neural net!



Fashion MNIST


- 70k Images
- 10 Categories
- Images are 28x28
- Can train a neural net!

<https://github.com/zalandoresearch/fashion-mnist>



```
import tensorflow as tf
from tensorflow import keras
```

```
fashion_mnist = keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```



Using a number is a first step in avoiding bias -- instead of labelling it with words in a specific language and excluding people who don't speak that language!


09

09 = ankle boot;
踝靴;
アンクルブーツ;
Bróg rúitín

<https://ai.google/responsibilities/responsible-ai-practices/>

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)
])
```

Flatten takes this 28 by 28 square and turns it into a simple linear array.



```

mnist = tf.keras.datasets.fashion_mnist
(training_images, training_labels), (test_images, test_labels) = mnist.load_data()
training_images=training_images/255.0 } Scale
test_images=test_images/255.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5)

```

Get Data (points to `mnist.load_data()`)

Scale (points to `255.0`)

Model (points to `tf.keras.models.Sequential`)

The training loop does support callbacks. So in every epoch, you can callback to a code function, having checked the metrics. If they're what you want to say, then you can cancel the training at that point.

```

class myCallback(tf.keras.callbacks.Callback):

```

```

    def on_epoch_end(self, epoch, logs={}):

```

```

        if(logs.get('loss')<0.4):

```

```

            print("\nLoss is low so cancelling training!")

```

```

            self.model.stop_training = True

```

sends logs
object about
current state
of training

When
Epoch
Ends

```

class myCallback(tf.keras.callbacks.Callback):

```

```

    def on_epoch_end(self, epoch, logs={}):

```

```

        if(logs.get('loss')<0.4):

```

```

            print("\nLoss is low so cancelling training!")

```

```

            self.model.stop_training = True

```

Query

← instantiate

```
callbacks = myCallback()
mnist = tf.keras.datasets.fashion_mnist
(training_images, training_labels), (test_images, test_labels) = mnist.load_data()
training_images=training_images/255.0
test_images=test_images/255.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5, callbacks=[callbacks])
```

Use Callbacks parameter ↗